

```
In [19]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns

#to ignore warnings
import warnings as wg
wg.filterwarnings("ignore")

In [21]: #Reading dataset from remote link
url="http://bit.ly/w-data"
data=pd.read_csv(url)
```

```
In [22]: print(data.shape)
data.head()
```

(25, 2)

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30

```
In [23]: data.describe()
```


Out[23]:

	Hours	Scores
count	25.000000	25.000000
mean	5.012000	51.480000
std	2.525094	25.286887
min	1.100000	17.000000
25%	2.700000	30.000000
50%	4.800000	47.000000
75%	7.400000	75.000000
max	9.200000	95.000000

```
In [24]: data.info()
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25 entries, 0 to 24
Data columns (total 2 columns):
Column Non-Null Count Dtype
--- -
0 Hours 25 non-null float64
1 Scores 25 non-null int64
dtypes: float64(1), int64(1)
memory usage: 528.0 bytes

```
In [25]: data.plot(kind='scatter',x='Hours',y='Scores');
plt.show()
```



```
In [26]: data.corr(method='pearson')
```

Out[26]:

	Hours	Scores
Hours	1.000000	0.976191
Scores	0.976191	1.000000

```
In [27]: data.corr(method='spearman')
```

Out[27]:

	Hours	Scores
Hours	1.000000	0.971891
Scores	0.971891	1.000000

```
In [28]: hours=data['Hours']
scores=data['Scores']
```

```
In [29]: sns.distplot(hours)
```



```
In [30]: sns.distplot(scores)
```




```
In [31]: x = data.iloc[:, :-1].values
y = data.iloc[:, 1].values
```

```
In [35]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2, random_state=50)
```

```
In [36]: from sklearn.linear_model import LinearRegression
reg=LinearRegression()
reg.fit(X_train, y_train)
```

```
Out[36]: LinearRegression()
```

```
In [37]: m=reg.coef_
c=reg.intercept_
line=m*X+c
plt.scatter(X,y)
plt.plot(X, line);
plt.show()
```




```
In [38]: y_pred=reg.predict(X_test)
```

```
In [39]: actual_predicted=pd.DataFrame({'Target':y_test,'Predicted':y_pred})
actual_predicted
```

Out[39]:

	Target	Predicted
0	95	88.211394
1	30	28.718453
2	76	69.020122
3	35	39.273652
4	17	13.365436

```
In [40]: sns.set_style('whitegrid')
sns.distplot(np.array(y_test-y_pred))
plt.show()
```



```
In [41]: h=9.25
s=reg.predict([[h]])
print("If a student studies for {} hours per day he/she will score {} % in exam.".format(h,s))

If a student studies for 9.25 hours per day he/she will score [91.56986604] % in exam.
```

```
In [45]: from sklearn import metrics
from sklearn.metrics import r2_score
print('Mean Absolute Error:',metrics.mean_absolute_error(y_test, y_pred))
print('R2 Score:',r2_score(y_test,y_pred))

Mean Absolute Error: 4.5916495300630285
R2 Score: 0.971014141329942
```

```
In [ ]:
```

```
In [ ]:
```